

Futures in TiKV

Peng Qu

2019.04.19

CONTENTS

What's futures and what's TiKV

Basic Future usage

Basic Stream usage

Collect Future into Stream

Executor-Task style batch system



What's Futures

What's Futures

- <https://github.com/rust-lang-nursery/futures-rs>

Traits including Future and Stream

mpsc/oneshot channels

Utilizes like executor/task

- <https://github.com/tokio-rs/tokio>

High level components in Future

tokio_threadpool, which can spawn Futures

```
use std::sync::mpsc as std_mpsc;

use futures::sync::mpsc;
use futures::{stream, Future, Sink, Stream};
use tokio_threadpool::Builder;

fn main() {
    let input = (0..4096).collect::<Vec<u32>>();
    let pool = Builder::new().pool_size(1).build();
    let (producer, consumer) = mpsc::channel(128);
    pool.spawn(
        stream::iter_ok::<_, ()>(input)
            .forward(producer.sink_map_err(|_| println!("sink error")))
            .map(|_| println!("produce finish")),
    );
    let (tx, rx) = std_mpsc::sync_channel(128);
    pool.spawn(
        consumer.for_each(move |i| Ok(tx.send(i).unwrap()))
            .map(|_| println!("consume finish")),
    );
    assert_eq!(rx.into_iter().count(), 4096);
    println!("finished");
}
```

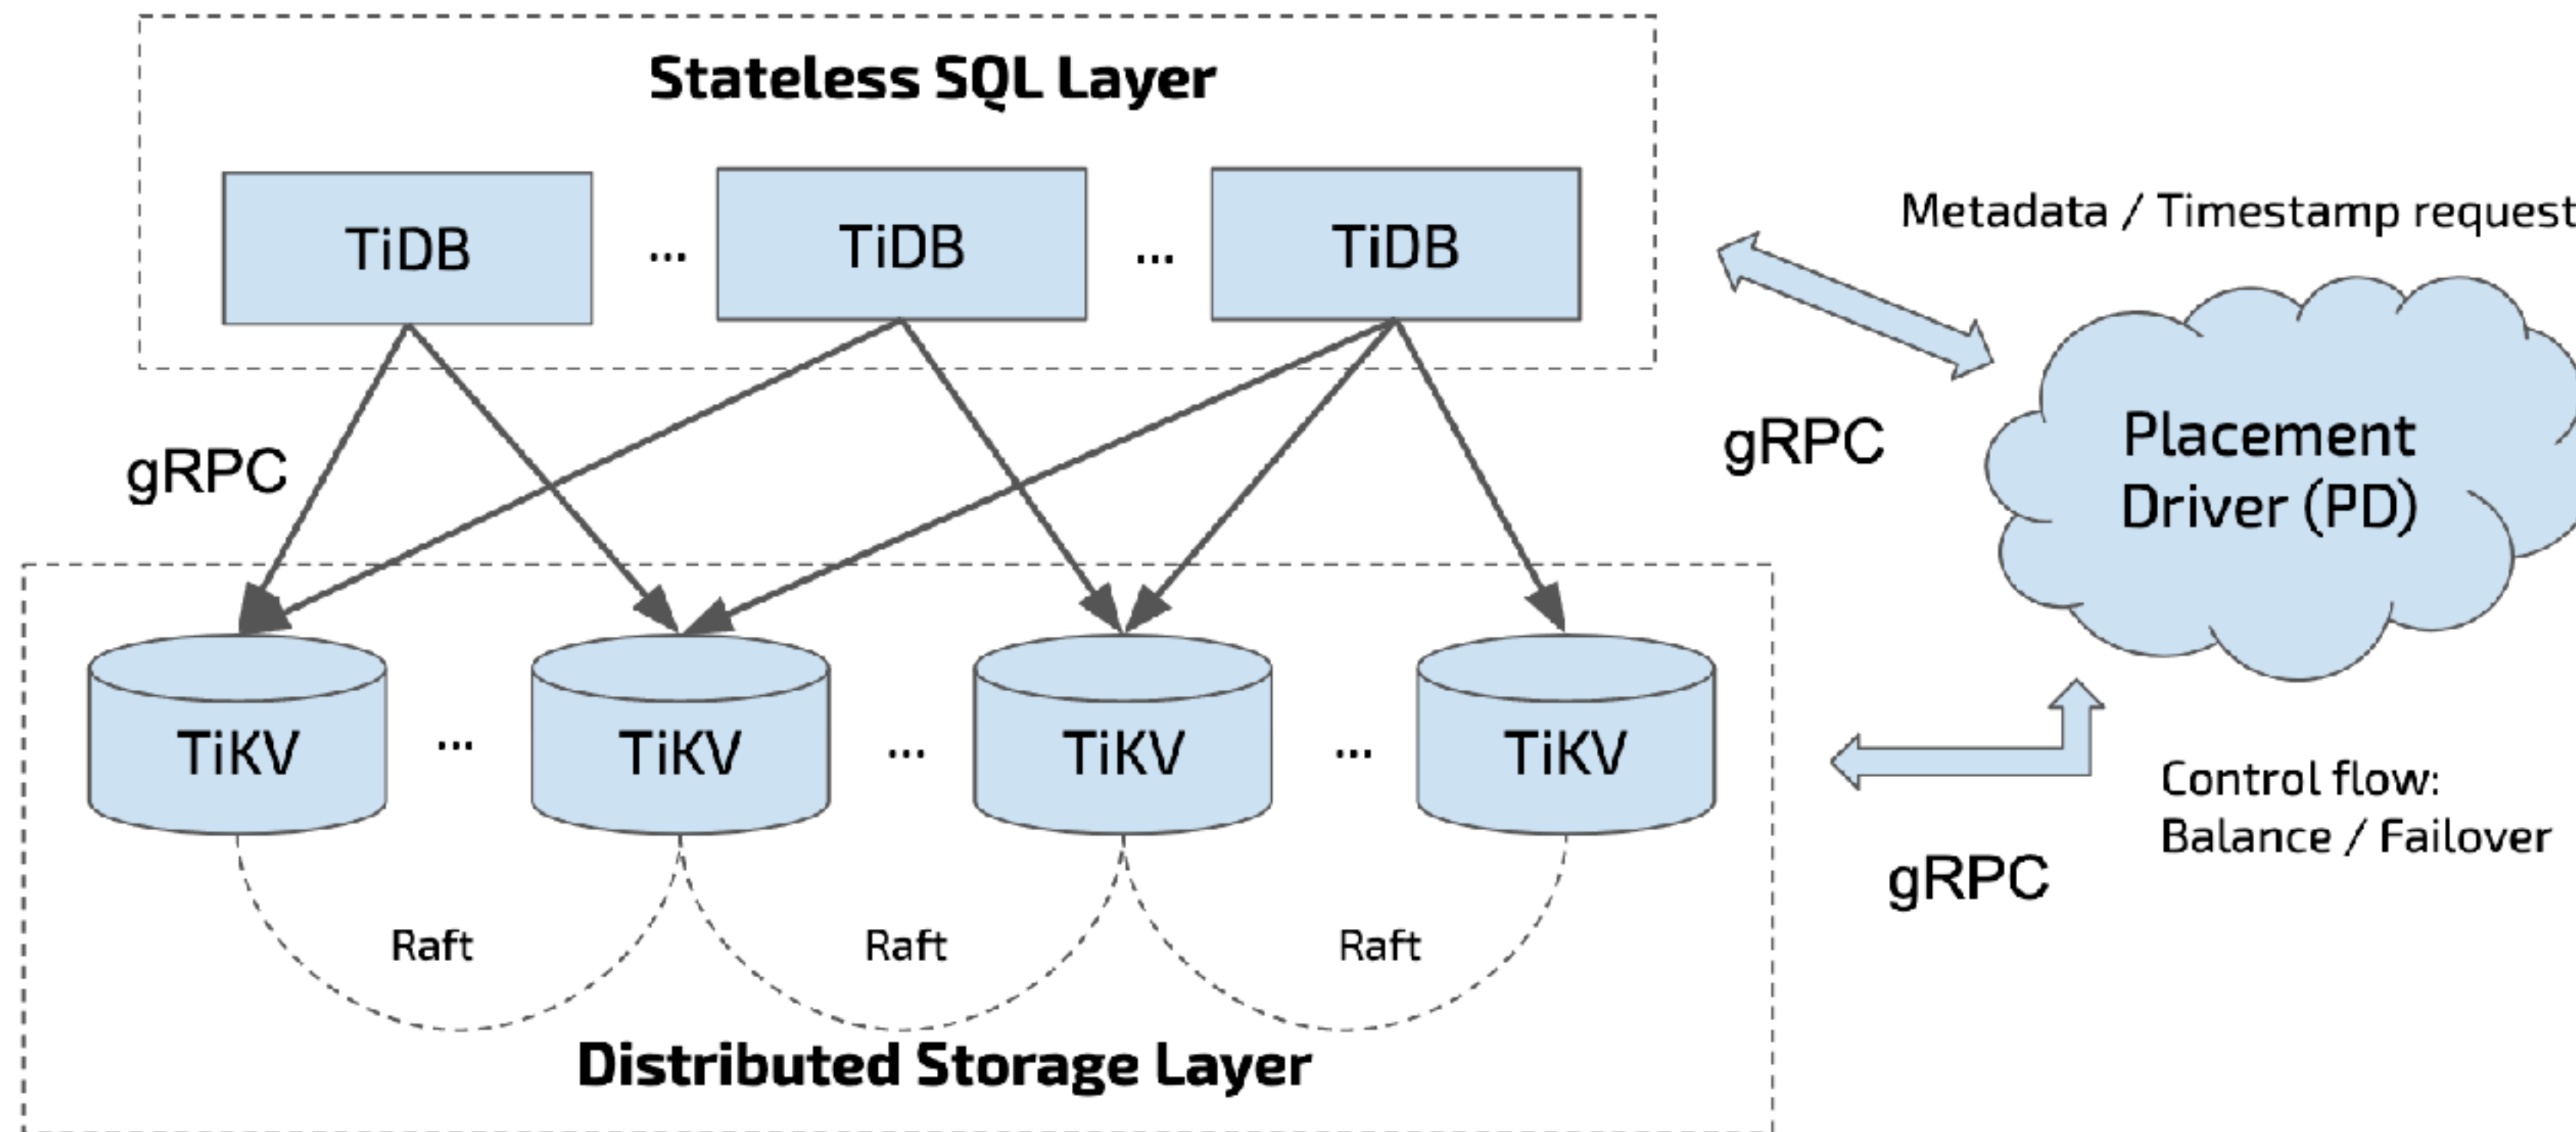


In the example, the producer and the consumer run in one thread

A large, red, starburst-shaped graphic with many sharp points, centered on a solid black background. The starburst has a slightly irregular, hand-drawn appearance.

What's TiKV

What's TiKV



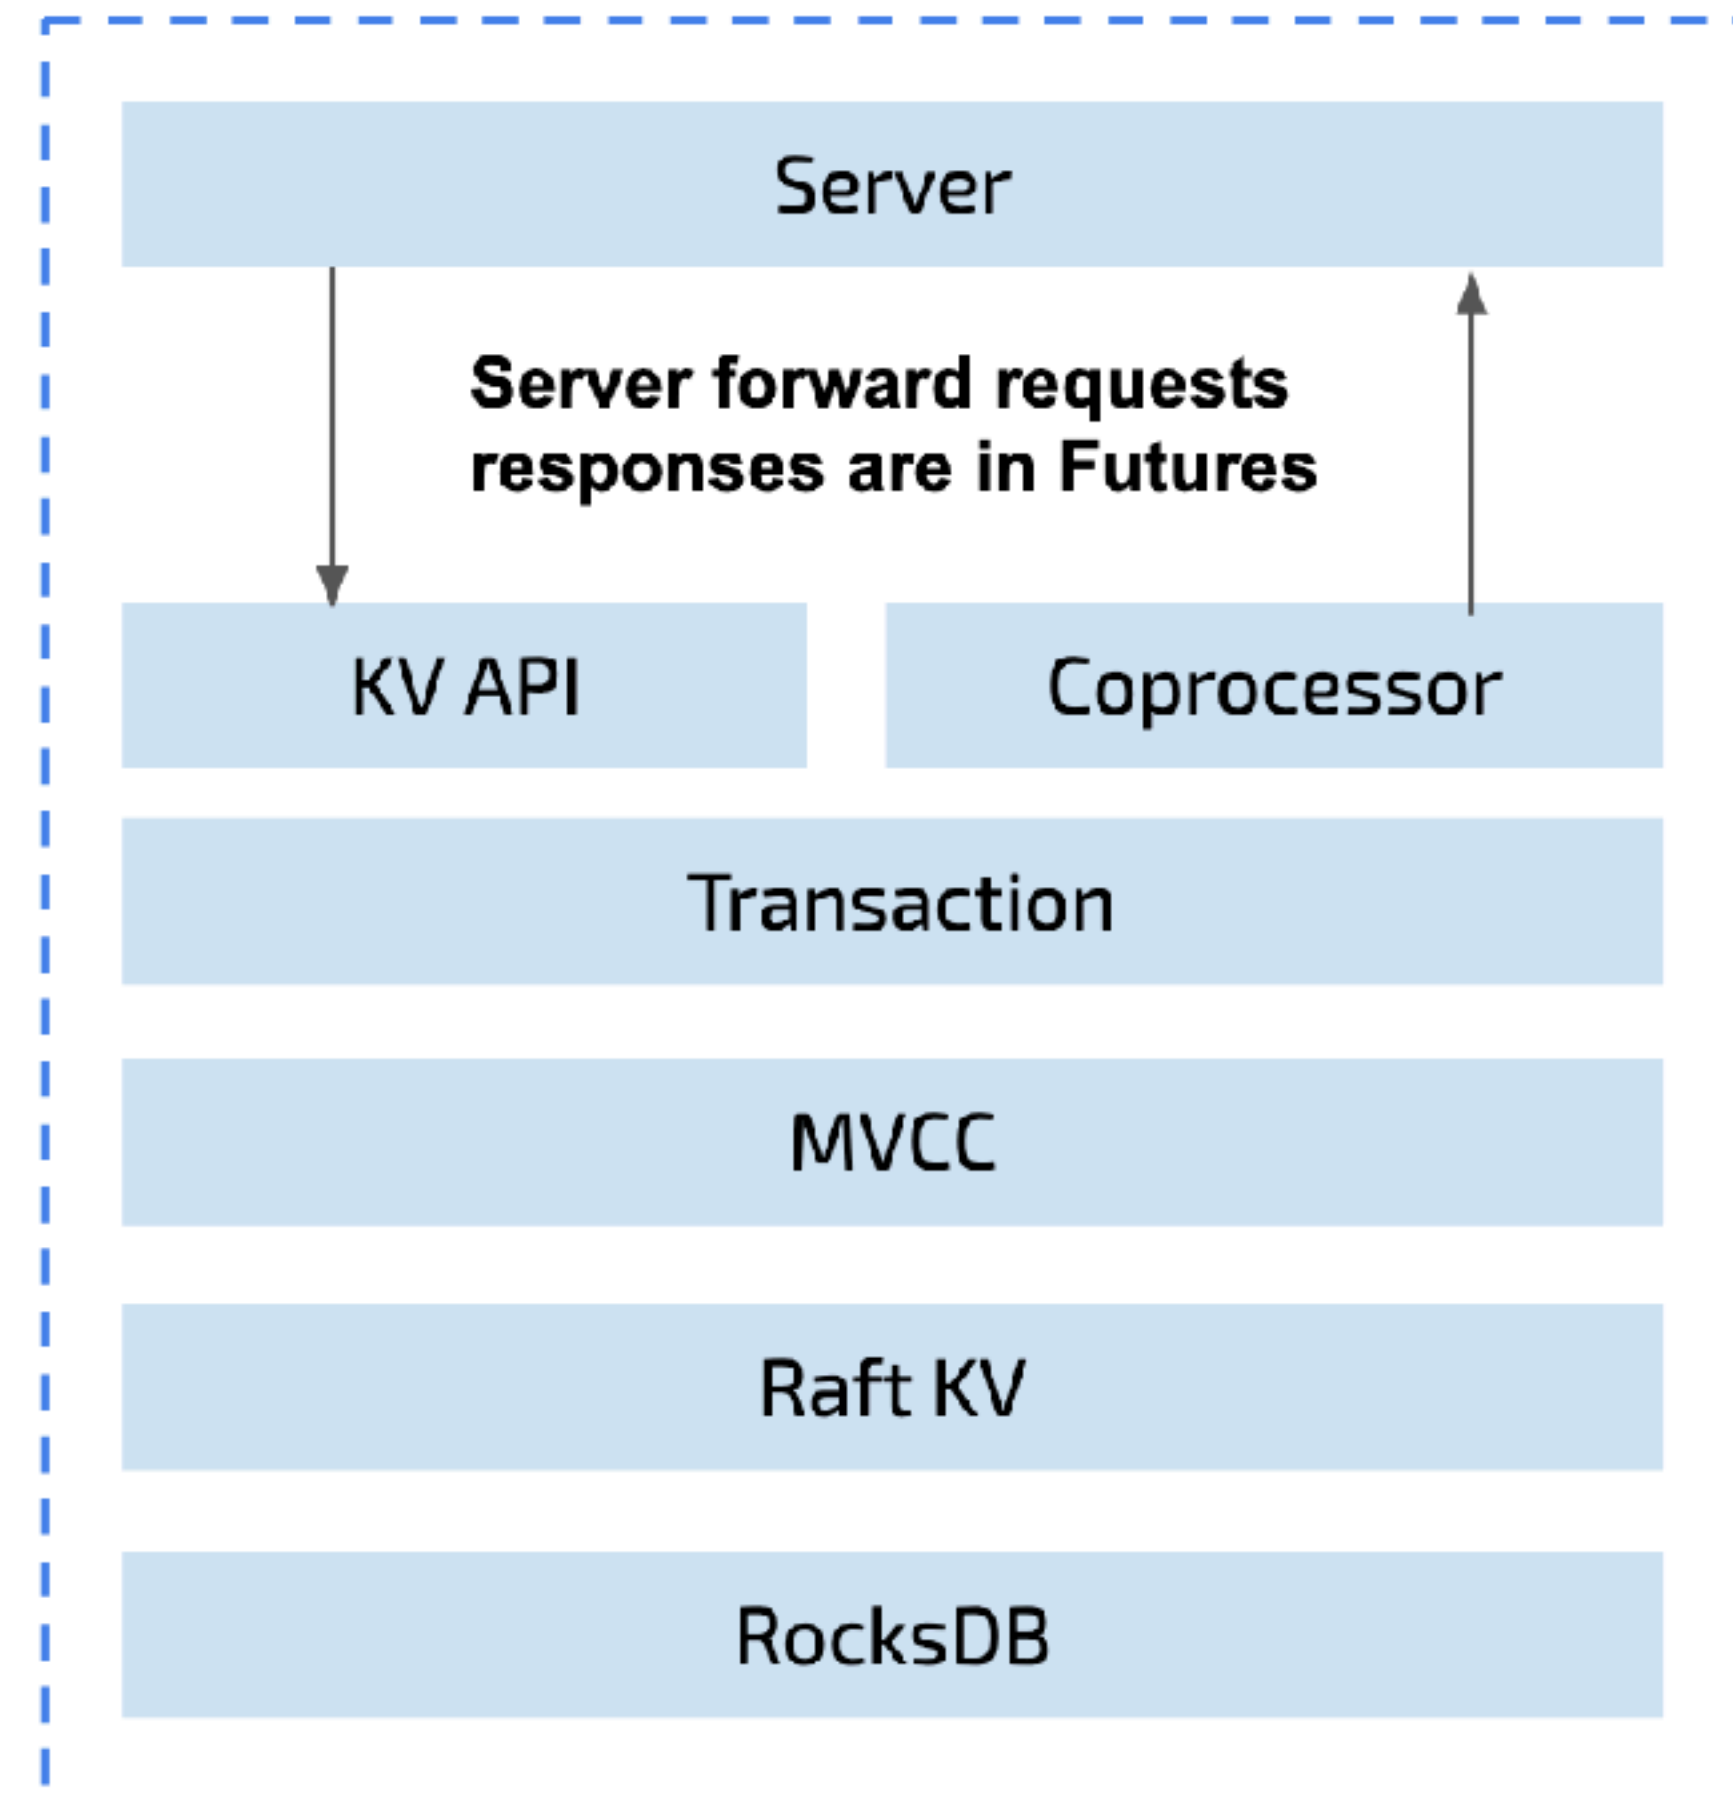
TiKV is a distributed KV storage engine. Data is replicated with Raft protocol

TiKV' internal & where are futures

Server layer retrieves requests, then forward them. Responses are in Futures, which are spawned in gRPC threads.

For some duplex streaming requests, there is also a usage in server layer that collects Futures into a Stream.

RaftKv appends users' write into Raft logs. There is a batch system inspired by Futures in it.



TiKV Layer Structure

A large, red, starburst-shaped graphic with many sharp points, centered on a black background. Inside the starburst, the text "Basic Future usage in TiKV" is written in a bold, black, sans-serif font.

Basic Future usage in TiKV

Basic Future usage in TiKV

```
impl tikvpb_grpc::Tikv for Service {  
    fn kv_get(&mut self, ctx: RpcContext<'_>, req: GetRequest, sink: UnarySink<GetResponse>) {  
        // Forward the request to background thread pool, and get a response future.  
        let f = future_get(&self.storage, req);  
        // Send the response back to the client.  
        let f = f.and_then(|res| sink.success(res).map_err(Error::from));  
        // Spawn the future into gRPC threads.  
        ctx.spawn(f);  
    }  
}
```

The key is how `spawn` is implemented

Basic Future usage in TiKV

```
impl<'a> RpcContext<'a> {  
    pub fn spawn<F>(&self, f: F:Future<Item = (), Error = ()> + Send + 'static>) {  
        // Spawn the future so that we can get its task.  
        let s = executor::spawn(Box::new(f) as BoxFuture<_, _>);  
        // Create a `Notify` which will kick the bounded completion queue if  
        // Notify::notify is called.  
        let notify = Arc::new(SpawnNotify(self.kicker));  
        // Pull the future in current thread.  
        poll(&notify, false);  
    }  
}
```

A large, red, starburst-shaped graphic with many sharp points, centered on a black background. It serves as a container for the title text.

Basic Stream usage in TiKV

Basic Stream usage in TiKV

```
impl tikvpb_grpc::Tikv for Service {  
    fn coprocessor_stream(  
        &mut self,  
        ctx: RpcContext<'_>,  
        req: Request,  
        sink: ServerStreamingSink<Response>,  
    ) {  
        let stream = self  
            .cop  
            .parse_and_handle_stream_request(req, Some(ctx.peer()))  
            .map(|resp| (resp, WriteFlags::default().buffer_hint(true)))  
            .map_err(|e| GrpcError::RpcFailure(RpcStatus::unknown));  
        ctx.spawn(sink.send_all(stream));  
    }  
}
```

A server streaming example, the key is the sink

Basic Stream usage in TiKV

```
impl tikvpb_grpc::Tikv for Service {  
    fn raft(  
        &mut self,  
        ctx: RpcContext<'_>,  
        stream: RequestStream<RaftMessage>,  
        sink: ClientStreamingSink<Done>,  
    ) {  
        let ch = self.ch.clone();  
        ctx.spawn(  
            stream.for_each(move |msg| ch.send_raft_msg(msg).map_err(Error::from))  
                .then(|res| {  
                    let status = match res {  
                        Err(e) => RpcStatus::new(RpcStatusCode::Unknown, Some(msg)),  
                        Ok(_) => RpcStatus::new(RpcStatusCode::Unknown, None),  
                    };  
                    sink.fail(status)  
                }  
            ),  
        );  
    }  
}
```

A client streaming example, the key is the stream

Basic Stream usage in TiKV (batch)

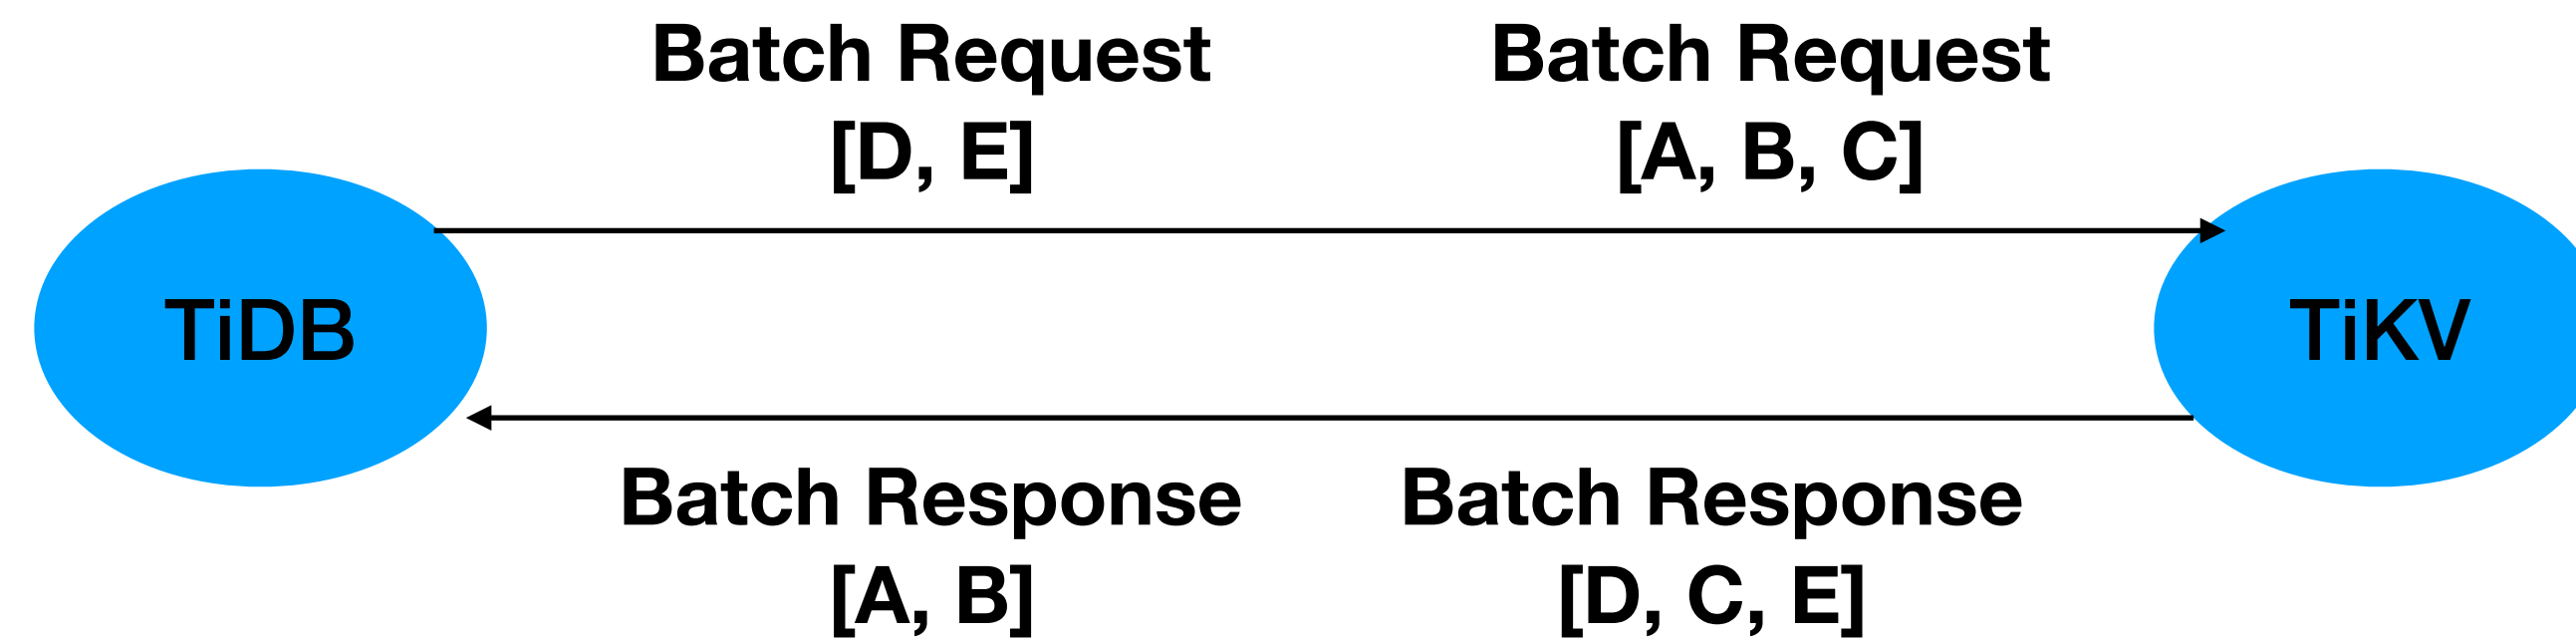
```
impl RaftClient {
    pub fn send(&mut self) { }
    pub fn flush(&mut self) {
        for conn in self.conns.values_mut() {
            // Get a notifier and do the notify explicitly.
            if let Some(notifier) = conn.stream.get_notifier() {
                if !self.grpc_thread_load.in_heavy_load() {
                    notifier.notify();
                    continue;
                }
                let wait = self.cfg.heavy_load_wait_duration.0;
                let _ = self.pool.spawn(
                    self.timer
                        .delay(Instant::now() + wait)
                        .inspect(move |_| notifier.notify()),
                );
            }
        }
    }
}
```

The BatchSender is in `tikv/components/tikv_util/src/mpsc/batch.rs`

A red starburst shape with many sharp points, centered on a black background. The text is written in black inside the starburst.

**Collect Futures
into Stream**

Collect Futures into Stream



- A batch request contains many little requests
- TiKV handles every little request in a Future
- TiKV needs to collect all Futures into a Stream

Collect Futures into Stream

```
fn collect_futures_into_stream(sink: Sender<Response>, f: impl Future<Item = Response, Error = ()>) {
    // Where to spawn `f`?
    let f = f.and_then(move |resp| sink.send(resp));
    // Don't need to spawn it into any threads, just set a callback.
    poll_future_notify(f);
}

fn poll_future_notify<F: Future<Item = (), Error = ()> + Send + 'static>(f: F) {
    let spawn = Arc::new(Mutex::new(Some(executor::spawn(f))));
    let notify = BatchCommandsNotify(spawn);
    notify.notify(0);
}

impl<F: Future<Item = (), Error = ()>> Notify for BatchCommandsNotify<F> {
    fn notify(&self, id: usize) {
        let n = Arc::new(self.clone());
        let mut s = self.0.lock().unwrap();
        match s.as_mut().map(|spawn| spawn.poll_future_notify(&n, id)) {
            Some(Ok(Async::NotReady)) | None => return,
            _ => *s = None,
        }
    }
}
```



A red starburst shape with many sharp points, centered on a black background. The text is written in black inside the starburst.

**Executor-Task
style batch system**

Executor-Task style batch system

- There are thousands regions per TiKV.
- Handle Raft messages could be a bottleneck.
- A mediocre implementation is easy to write, but
 - some regions could be hungry
 - hard to consolidate a write batch
 - hard to refactor old code
- So we need a better implementation...

```
use std::sync::mpsc as std_mpsc;

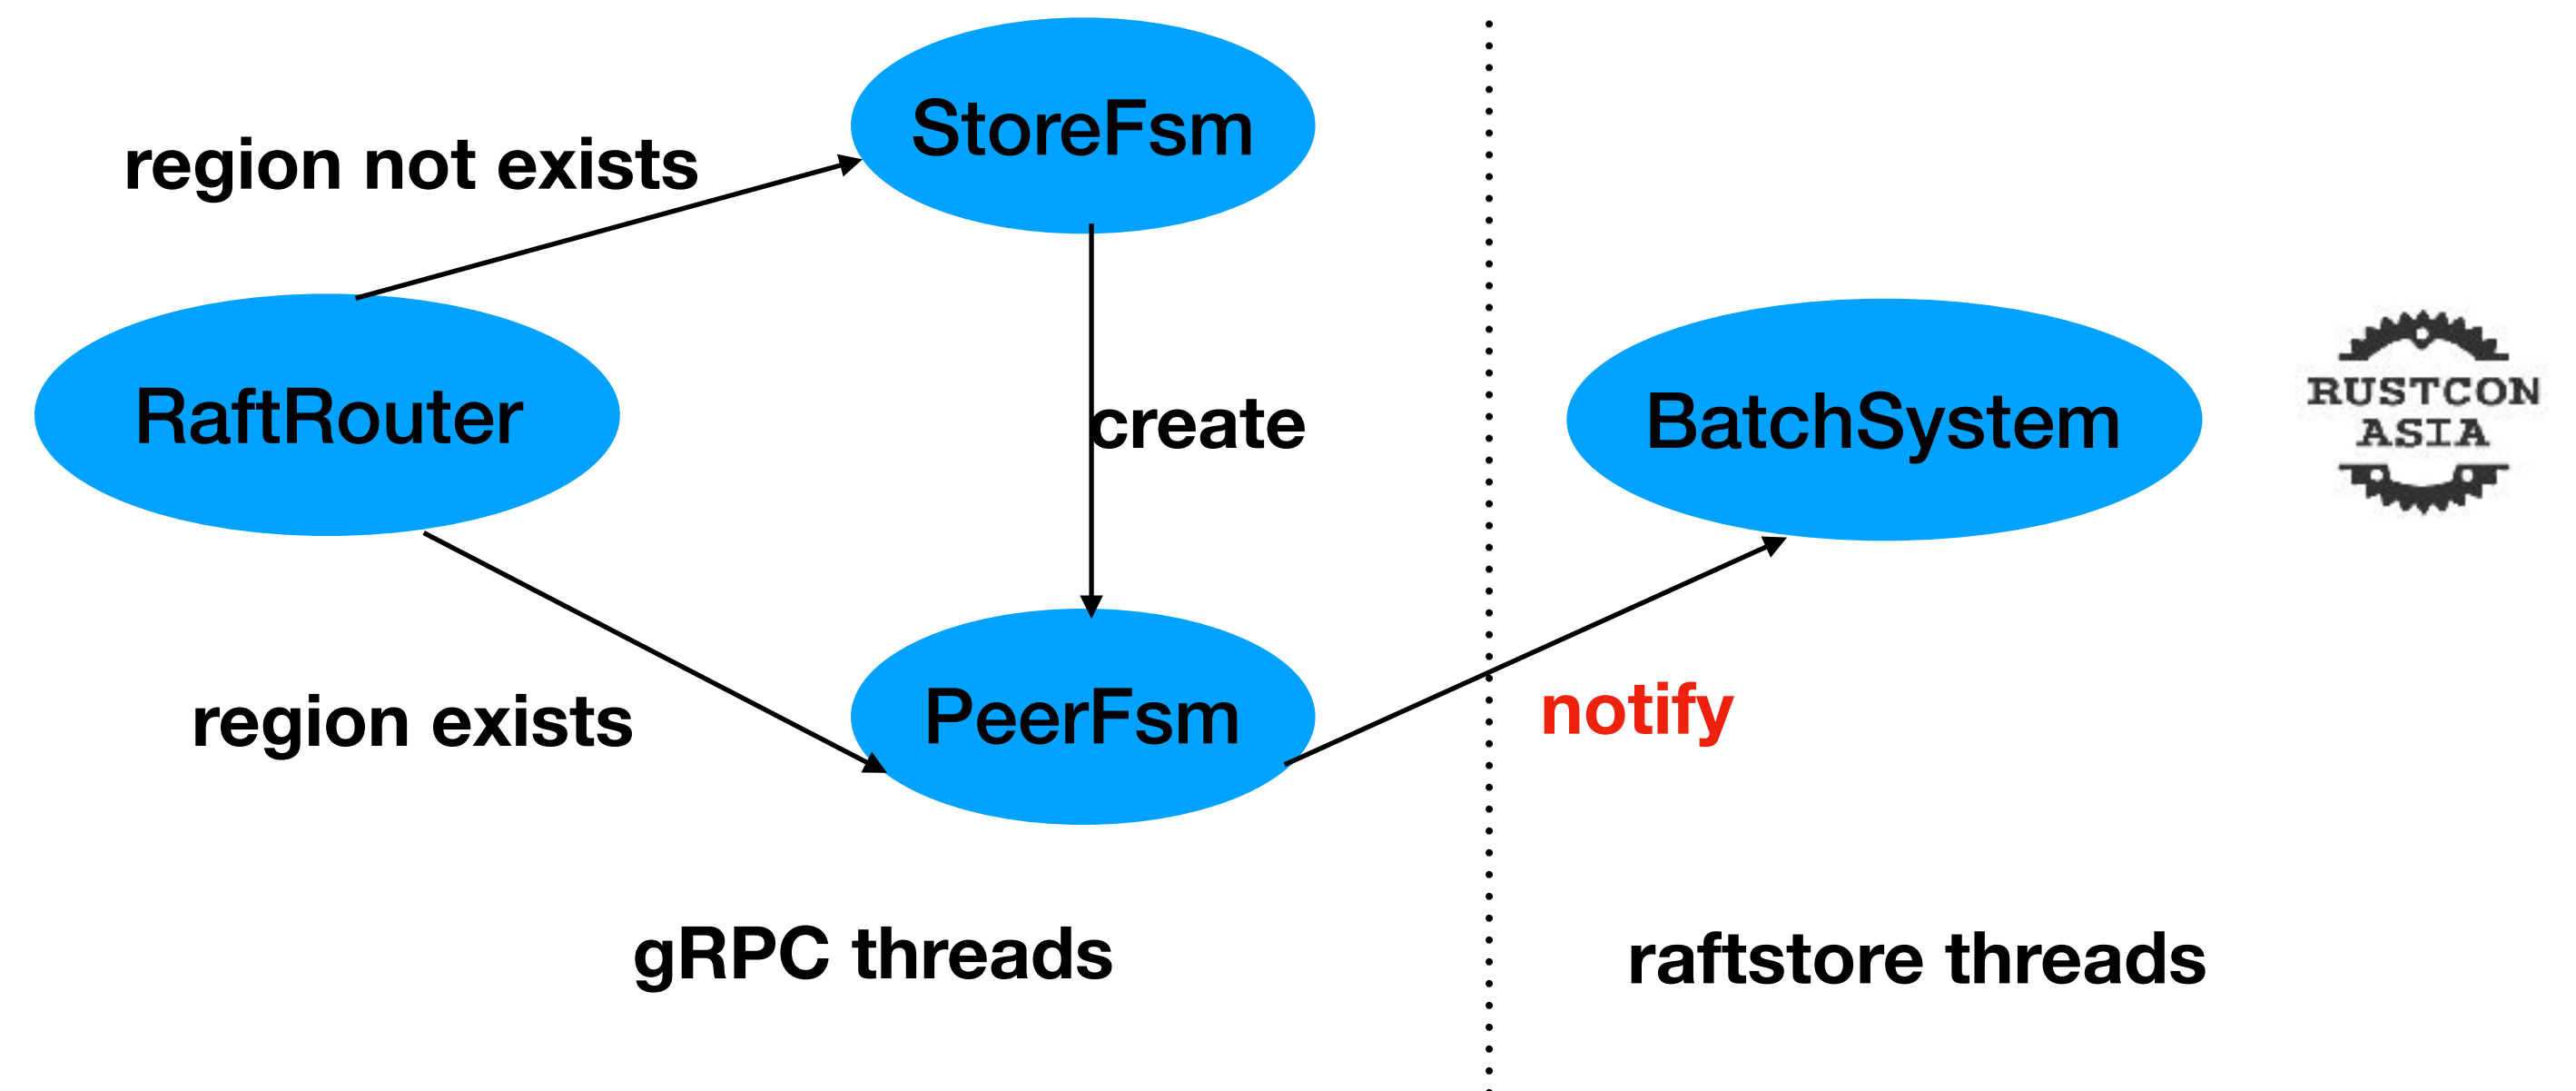
use futures::sync::mpsc;
use futures::{stream, Future, Sink, Stream};
use tokio_threadpool::Builder;

fn main() {
    // Suppose we have 4 threads to handle Raft messages.
    let executor = Builder::new().pool_size(4).build();
    // Suppose we have 10k regions on a TiKV.
    let mut regions = Vec::with_capacity(10000);
    for _ in 0..10000 {
        let (_, rx) = mpsc::channel(1024);
        regions.push(rx);
        executor.spawn(rx.for_each(|msg| handle_raft_msg(msg)));
    }
}
```



Executor-Task style batch system

- Fsm means Finite-State Machine.
- PeerFsm will notify BatchSystem if need.
- In BatchSystem:
 - many background threads
 - batch strategy to avoid hungry
 - hung a region if it's temporarily unavailable



Executor-Task style batch system

```
// RaftRouter will call `MailBox::send` and `MailBox::notify` to forward
// Raft messages and notify downstream components to handle them.
impl BasicMailBox {
    // FsmScheduler can schedule a PeerFsm to a BatchSystem to run.
    fn notify(&self, scheduler: &FsmScheduler) {
        let prev_state = self.state.compare_and_swap(IDLE, NOTIFIED);
        if prev_state == IDLE {
            // IDLE means the PeerFsm is not running. So schedule it.
            let mut owner = unsafe { Box::from_raw(self.data) };
            owner.set_mailbox(Cow::Borrowed(self));
            scheduler.schedule(owner);
        }
    }
    // Called in BatchSystem to put `fsm` back to the mailbox.
    fn release(&self, fsm: Box<Owner>) {
        assert!(self.data.swap(Box::into_raw(fsm)).is_null());
        let prev_state = self.state.compare_and_swap(NOTIFIED, IDLE);
        if prev == DROP {
            // It's destroyed instead of released, deconstruct it.
            drop(unsafe { Box::from_raw(self.data) });
        }
    }
}
```

Implementation of notify and release




Executor-Task style batch system

- A batch contains many regions, so no hungry.
- It's easy to consolidate a write batch from many regions.
- More easily to be integrated into old code

```
impl BatchSystemPoller {  
    // The loop of BatchSystem's background threads.  
    fn poll(&mut self) {  
        // `Batch` is like a Vector of Owner.  
        let mut batch = Batch::with_capacity(self.max_batch_size);  
        self.fetch_batch(&mut batch);  
        while !batch.is_empty() { // It's empty means the region is destroyed.  
            let mut raft_readies = Vec::new();  
            for peer in batch.iter() {  
                raft_readies.push(self.handle_raft_messages(&peer));  
                if peer.temporarily_unavailable() {  
                    batch.release(peer);  
                }  
            }  
            self.handle_raft_readies(raft_readies);  
            self.fetch_batch(&mut batch);  
        }  
    }  
}
```

Implementation of BatchSystem's loop





THANKS